

DC-URS v1 — implementation plan

Contract: `dc-urs-spec.md` Date: 2026-07-27 Estimate: ~3 days across five phases, each landing green before the next.

Every phase ends with `npm run verify` passing and a commit. Phase 0 is independent of all data acquisition and can start immediately.

Audit findings — read before planning anything

The calibration plan was audited against the code before execution and four assumptions turned out wrong. Same exercise here; two findings change the plan materially.

The replacement surface is much smaller than expected

`greenScore` has **exactly one call site** — `heat-map-app.ts:350` — and only three files mention it at all: the model, the app, and the Astro shell. No tests, no compare page, no scenario URL encoding. The swap is genuinely localised.

The UI surface is also close to a 1:1 match. The existing `.subs` strip shows **three raw sub-scores** because a composite index is only auditable if you can see which component produced the number. DC-URS has **three pillars**. The strip becomes ACI / EVI / THI with almost no layout work.

DC-URS has no economics, and that is a real product regression

The Green Score's third component is **budget efficiency** — °C per crore. That is why trees score 84 and facades score 4: the index rewards cost-effectiveness, and that ordering matches Indian Heat Action Plan practice.

DC-URS contains no cost term at all. `ACI = UGS + CRI + TRA` ; nothing in the index knows what anything costs. So on replacement:

- a ward that achieves its greenery cheaply and one that spends ₹200 crore to get there **score identically**
- the "which intervention is best value" answer — arguably the most decision-useful thing the tool produces — disappears from the score

This is not a defect in DC-URS. It is a *state* index, not an *efficiency* index, and it is correct for that job. But the capability is lost unless we deliberately keep it. **Flagged to the CEO as a decision** (§Open questions); the plan below keeps cost computed and displayed so the option stays open, and adds nothing to the index without his sign-off.

Smaller findings

- `WARDS` is **inline in** `heat-map-app.ts`, not in `src/data/`. Spec §1 requires it to be data. Extraction is Phase 0 work, not a refactor to do later.
- `state.greenG` (`computeGreenG`) feeds the Green Score's greening ratio. DC-URS does not use it, but it is closely related to `FVC` and should **not** be deleted until Phase 4 confirms nothing else reads it.
- `scoreTxt` currently renders cooling, greening and cost together. It needs rewriting regardless of the cost decision.

File plan

<code>src/data/</code>	
<code>wards.ts</code>	NEW P0 · ward list extracted from <code>heat-map-app.ts</code>
<code>src/scripts/climate-engine/</code>	
<code>dc-urs.ts</code>	NEW P0 · the engine, pure and self-checked
<code>dc-urs-inputs.ts</code>	NEW P0 · typed indicator record + per-field provenance
<code>heat-map-app.ts</code>	EDIT P4 · swap the call site, wire the pillars
<code>heat-map-model.ts</code>	EDIT P4 · retire <code>greenScore/greenScoreParts</code>
<code>src/components/ClimateEngine/</code>	
<code>HeatMapStage.astro</code>	EDIT P4 · sub-scores → pillars, structural floor, tier label
<code>scripts/</code>	
<code>fetch-worldpop.py</code>	NEW P2 · population density per footprint
<code>fetch-census-2011.py</code>	NEW P2 · HVI_socio, areal-interpolated
<code>fetch-sentinel-composites.py</code>	NEW P1 · seasonal FVC, albedo, multi-year NDVI series
<code>compute-far.py</code>	NEW P1 · FAR from Google 2.5D + MS footprints
<code>compute-tra.py</code>	NEW P1 · distance-to-refuge from OSM
<code>build-dcurs-inputs.py</code>	NEW P1 · assemble → <code>data/dc-urs/inputs.json</code>
<code>measure-dcurs-uncertainty.py</code>	NEW P3 · Monte Carlo band
<code>validate-dcurs.mjs</code>	NEW P3 · the four spec §7 checks
<code>data/dc-urs/</code>	
<code>inputs.json</code>	NEW P1 · one record per ward, every field with provenance
<code>golden-cases.json</code>	NEW P0 · frozen regression fixture

`dc-urs.ts` is separate from `heat-map-model.ts` because it is a different kind of thing: the model computes physics, the index composes indicators. Mixing them is what made the Green Score hard to audit.

Phase 0 — The engine (½ day, unblocked)

No data acquisition. Everything here runs on hand-fed inputs.

0a • `src/data/wards.ts`

- [] Move the `WARDS` object out of `heat-map-app.ts` verbatim, add the fields DC-URS needs.

```
export interface Ward {
  readonly id: string;
  readonly name: string;
  readonly zone: string;
  readonly body: string;           // KMC Ward 68 | Baruipur Municipality | ...
  readonly lat: number;
  readonly lon: number;
  readonly footprintM: number;    // 1400 – matches the thermal simulation domain
}

export const WARDS: readonly Ward[] = [ /* ... */ ];
```

- [] `heat-map-app.ts` imports from here. **Adding a ward must never require touching a script.**

0b • `src/scripts/climate-engine/dc-urs-inputs.ts`

- [] The eleven indicator values with provenance per field, reusing the `Sourced<T>` idea already established in the project.

```
export type Provenance = 'measured' | 'modelled' | 'estimated' | 'placeholder';
export interface Sourced<T> { value: T; source: Provenance; vintage?: string; cite?: string }

export interface DcUrsInputs {
  lstDayC: Sourced<number>;   lstNightC: Sourced<number>;  ruralBaseC: Sourced<number>;
  fvc: Sourced<number>;       canopyFrac: Sourced<number>;
  ndviMean: Sourced<number>;  ndviStd: Sourced<number>;
  albedo: Sourced<number>;    distCoolM: Sourced<number>;
  popDensity: Sourced<number>; far: Sourced<number>;  socioVuln: Sourced<number>;
}
```

Provenance is per field because it lands field by field, and the UI must mark what is 2011 census versus what is this morning's satellite pass.

0c • `src/scripts/climate-engine/dc-urs.ts`

- [] Anchors as named constants, each carrying its provenance in a comment (spec §4).
- [] The corrected engine, exactly as spec §2:

```
/** Vegetation stability. Guarded: water's negative mean NDVI made the unguarded
 * form return 1.00 – perfect vegetation stability for a river. */
export function vsi(ndviMean: number, ndviStd: number): number {
  if (ndviMean <= VSI_NDVI_FLOOR) return 0;          // 0.10
  return 1 - clamp(ndviStd / ndviMean, 0, 1);
}

export function pillars(i: DcUrsInputs): { aci: number; evi: number; thi: number; ugs:
number } {
  const thi = 0.40 * clamp((i.lstDayC.value - 25) / 25, 0, 1)      // /25, not /20 –
delta 7
    + 0.40 * clamp((i.lstNightC.value - 20) / 15, 0, 1)
    + 0.20 * clamp(Math.max(0, i.lstDayC.value - i.ruralBaseC.value) / 10, 0, 1);
  const evi = 0.45 * clamp(i.popDensity.value / 25000, 0, 1)
    + 0.30 * clamp(i.far.value / 5, 0, 1)
    + 0.25 * clamp(i.socioVuln.value / 10, 0, 1);
  const ugs = 0.50 * clamp(i.fvc.value, 0, 1)
    + 0.30 * clamp(i.canopyFrac.value, 0, 1)
    + 0.20 * vsi(i.ndviMean.value, i.ndviStd.value);
  const aci = 0.50 * ugs + 0.30 * clamp(i.albedo.value / 0.60, 0, 1)
    + 0.20 * Math.exp(-0.002 * i.distCoolM.value);
  return { aci, evi, thi, ugs };
}

/** Geometric, not additive. IPCC AR6 risk is conjunctive: a near-zero pillar must
 * drag the score down rather than be bought off by a strong one. */
export function dcUrs(i: DcUrsInputs): number {
  const { aci, evi, thi } = pillars(i);
  const g = (x: number) => clamp(x, 0.001, 1);          // never evaluate 0^0.4
  return 100 * g(aci) ** 0.40 * g(1 - evi) ** 0.35 * g(1 - thi) ** 0.25;
}
```

- [] `tierFor(score)` returning the spec §6 band, label and colour.
- [] `structuralFloor(i)` — a **named export**, not an incidental number:

```

/** Points that no intervention can recover, because EVI responds to nothing the
 * user can do. This is the tool's sharpest output, not a limitation: it separates
 * what greening can fix from what only housing and health policy can. */
export function structuralFloor(i: DcUrsInputs): { ceiling: number; withheld: number } {
  const perfect: DcUrsInputs = { ...i,
    fvc: { ...i.fvc, value: 1 }, canopyFrac: { ...i.canopyFrac, value: 1 },
    ndviMean: { ...i.ndviMean, value: 0.85 }, ndviStd: { ...i.ndviStd, value: 0.03 },
    albedo: { ...i.albedo, value: 0.60 }, distCoolM: { ...i.distCoolM, value: 0 },
    lstDayC: { ...i.lstDayC, value: i.ruralBaseC.value },
    lstNightC: { ...i.lstNightC, value: i.lstNightC.value - 5 },
  };
  const ceiling = dcUrs(perfect);
  return { ceiling, withheld: 100 - ceiling };
}

```

0d · Self-check and golden cases

- [] `assertDcUrsLogic()` — runnable with `node --experimental-strip-types`:

```

export function assertDcUrsLogic(): void {
  const a = (ok: boolean, m: string) => { if (!ok) throw new Error(`dc-urs: ${m}`); };
  a(vsi(-0.30, 0.05) === 0, 'water must score 0 stability, not 1');
  a(vsi(0.05, 0.06) === 0, 'bare ground must score 0 stability');
  a(vsi(0.75, 0.04) > 0.9, 'mature canopy must score high stability');
  // bounded, and monotone in each pillar
  for (const w of GOLDEN) {
    const s = dcUrs(w.inputs);
    a(s >= 0 && s <= 100, `${w.id}: score out of range`);
    a(Math.abs(s - w.expected) < 0.05, `${w.id}: ${s.toFixed(2)} vs frozen
    ${w.expected}`);
  }
  // a lethal pillar cannot be bought off – the whole point of geometric aggregation
  a(dcUrs(HOT_BUT_GREEN) < dcUrs(COOL_BUT_BARE), 'hot ward must not outscore cool ward');
  a(structuralFloor(BALLYGUNGE).ceiling < 100, 'structural floor must withhold
  something');
}

```

- [] Freeze `data/dc-urs/golden-cases.json` from the dry run already performed: **Ballygunge 21.13**, **Baruipur 64.16**. Barrackpore is *excluded* until Phase 1 supplies real inputs — its dry-run 40.14 used estimated values and must not be frozen as truth.

Verify: `assertDcUrsLogic()` passes; `npm run check` 0 errors. Nothing user-visible changes yet.

Phase 1 — Observed pillars (1 day)

Real satellite and vector data for everything except demographics.

- [] `fetch-sentinel-composites.py` — Sentinel-2 L2A via the existing EE service account. **Seasonal composites, never single scenes** (spec §3): Kolkata NDVI swings hard between monsoon and dry season. Emit per ward: `fvc`, `albedo`, and a ≥ 5 -year annual NDVI series for `ndviMean` / `ndviStd`. Albedo uses the source document's coefficients — note in the header that this is a Liang-type narrowband-to-broadband conversion so the lineage is traceable.
- [] `compute-far.py` — floor-area ratio from Google Open Buildings 2.5D heights over Microsoft footprints, both already held. $FAR = \Sigma(\text{footprint area} \times \text{floors}) / \text{footprint area}$, floors estimated at 3.2 m storey height, assumption stated in the output.
- [] `compute-tra.py` — nearest thermal refuge from OSM (`leisure=park`, `natural=water`, `amenity=community_centre`), distance transform over the ward footprint, mean distance.
- [] `canopyFrac` — ESA WorldCover class 10 fraction, reusing `suhii.align()` and the 70 m target grid so it lands on the same pixels as the thermal work.
- [] `build-dcurs-inputs.py` — assemble into `data/dc-urs/inputs.json`, one record per ward, **every field carrying its provenance and vintage**.

Verify: every observed field populated for all three wards; no field left `placeholder`; run the engine on real inputs and record the scores. Expect movement from the dry run — the dry run used the CEO's illustrative numbers.

Watch for: `canopyFrac` from a top-of-canopy product suppresses built fraction in dense tropical cities (street trees occlude roofs). Record the caveat; do not correct for it silently.

Phase 2 — Demographics (½ day)

- [] `fetch-worldpop.py` — WorldPop constrained, current release, clipped to each 1400 m footprint → `popDensity` in persons/km². CC BY 4.0, attribution recorded.
- [] `fetch-census-2011.py` — `HVI_socio` from Census of India 2011, composed the Rath (2021) / Azhar (2017) way: elderly >65, children <5, low-income proxy, informal settlement fraction. **Areal interpolation** onto the 1400 m footprint, because the three wards sit under three different local bodies and no single body's ward figures cover them (spec §1).
- [] Vintage surfaced in `inputs.json` and rendered in the UI. **Census 2021 does not exist** — postponed to Census 2027, reference date 1 March 2027.

Verify: EVI complete for all three wards; population sanity-checked against published Kolkata ward densities; the WorldPop-vs-Census gap recorded rather than reconciled.

Phase 3 — Validation (½ day)

`validate-dcurs.mjs`, the four spec §7 checks. Rank correlation is **not** among them — it is impossible at $n=3$ and that is stated, not worked around.

- [] **Golden-case regression** — the frozen fixture must reproduce exactly.
- [] **Face validity** — ranking must match known Kolkata geography. A dense hot core outscoring a green fringe is a failure whatever the maths says.
- [] **Sensitivity analysis** — vary each of the eleven inputs across its plausible range, record the DC-URS swing, print a ranked table. **If one indicator dominates, the other ten are decoration and the weighting is wrong.** Report; do not silently retune.
- [] `measure-dcurs-uncertainty.py` — Monte Carlo the thermal model's measured $\pm 3.5^\circ\text{C}$ night / $\pm 5.0^\circ\text{C}$ day through the geometric form. The additive-era estimate of ~ 5 points **does not carry over** and must be re-measured (spec §6).

Verify: all four run and report. No numeric bar is set for face validity or sensitivity — inventing one at $n=3$ would repeat the $\pm 2\text{ K}$ mistake from the thermal calibration, where a bar was set before anyone knew what was achievable.

Phase 4 — Wire it in, retire the Green Score (½ day)

- [] Replace the single call site at `heat-map-app.ts:350`. Baseline pillars come from `inputs.json`; scenario pillars apply modelled deltas to `fvc`, `canopyFrac`, `albedo`, `distCoolM` and the LST terms. **vsi is frozen at baseline** — newly planted cover scores ≈ 0.27 against mature canopy's 0.95, so a live VSI would make planting trees *lower* the score.
- [] `.subs` strip: greening/cooling/efficiency $\rightarrow$ **ACI / EVI / THI**, raw, same auditability rationale as before.
- [] Tier label and colour from `tierFor()`, including "Critical Hotspot" — CEO-approved wording.
- [] **Structural floor rendered**, not hidden: achievable headroom shown separately from withheld points, with the sentence that makes it useful — *"28 of this ward's missing points are demographic exposure; trees cannot fix them."*
- [] Uncertainty band beside the score, matching the treatment already shipped for temperature. Baseline uses observed temperature; scenario carries the model band. **The asymmetry is visible.**
- [] Delete `greenScore` / `greenScoreParts` from `heat-map-model.ts`. Keep `computeCost` and `computeGreenG` until the cost decision is made.
- [] Update `docs/green-score-methodology.md` — the index it documents no longer ships.

Verify: `npm run verify green`; e2e passes; sliders move DC-URS; no console errors; the score never leaves 0–100.

Open questions for the CEO

1. **Does the tool keep a cost-effectiveness signal?** DC-URS has no economics, so trees and green facades that achieve the same greenery score identically regardless of costing ₹66 lakh or ₹190 crore. Options: (a) accept the loss, cost shown as context only; (b) display ₹/point alongside the score without touching the index; (c) add a fourth pillar — a change to his engine. **Plan assumes (b)** as the reversible default.
 2. **What replaces `docs/green-score-methodology.md`?** The consultant reviewed it eight days ago. A short "what changed and why" note is probably owed.
-

What could go wrong

	Symptom	Response
Golden cases drift	Frozen values don't reproduce	A real regression. Do not re-freeze to make it pass
Sensitivity dominated by one input	One indicator moves the score, ten don't	Weighting problem. Report with the table; CEO decides
Seasonal NDVI ambiguity	FVC differs wildly by composite window	Fix the window in the spec and record it; do not average monsoon with dry
Census interpolation implausible	A ward's density is obviously wrong	Say so. Do not smooth it into looking reasonable
Wards land in one tier	Real inputs compress the range	Report before adjusting anchors — the anchors are the CEO's and researched
Score jumps vs Green Score	Users see a different number	Expected and pre-agreed. Document the delta per ward
Uncertainty band is large	± 10 points or worse	Publish it. That is the honest outcome and the precedent is set

Sequencing rule

One change at a time, measured after each. The failure this project has already seen twice — a constant tuned until the picture looked right — is prevented by the golden-case fixture and the sensitivity table, not by good intentions. Every phase ends with a number that either matches expectation or does not.